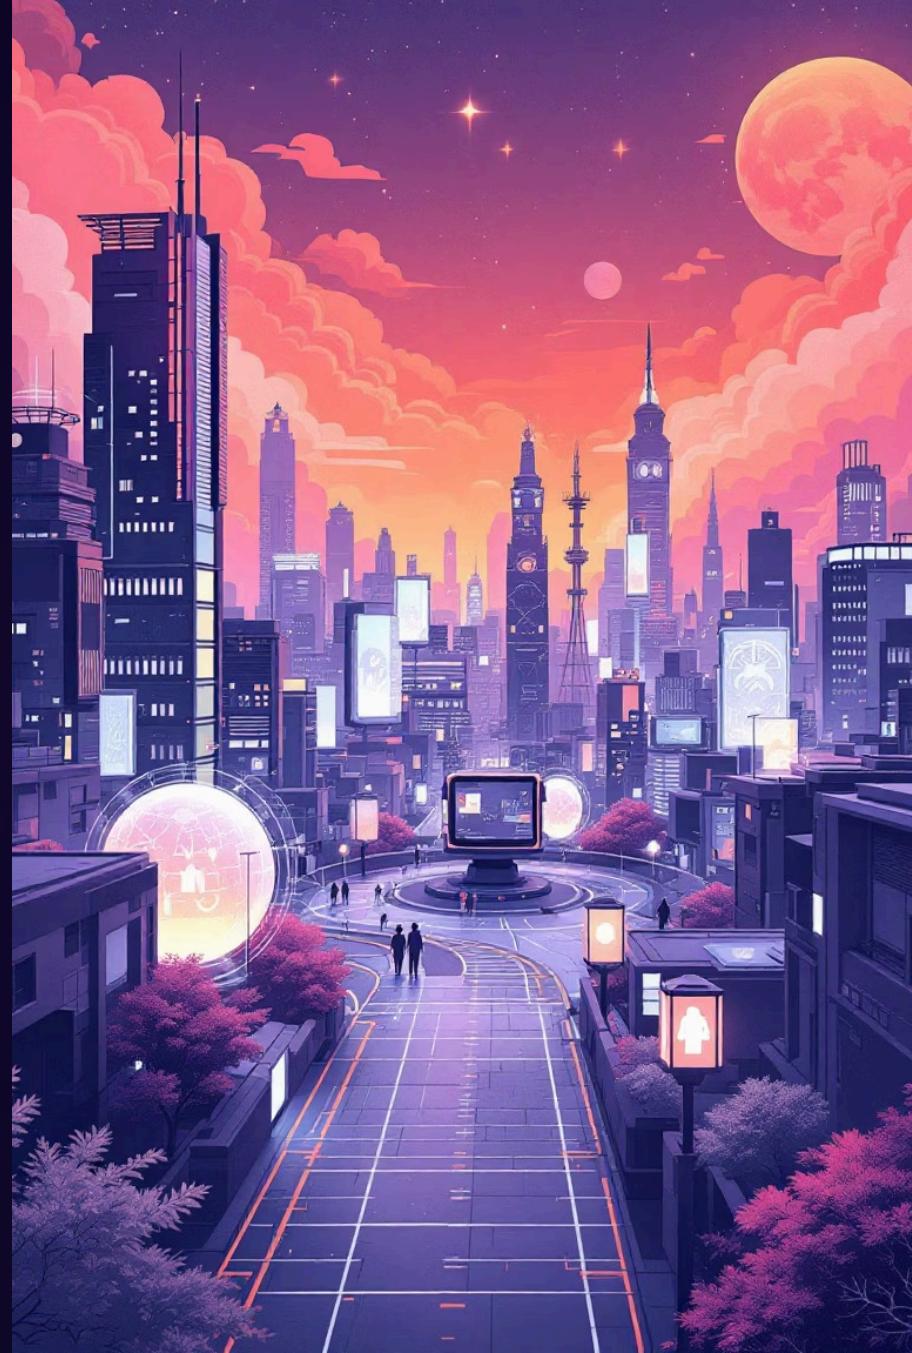


# Artificial Intelligence: Understanding Environments and Knowledge-Based Agents

Welcome to Week 4 of our Artificial Intelligence lecture series! Today, we'll dive into how AI agents perceive and interact with their surroundings, exploring the crucial concepts of environments and the powerful idea of Knowledge-Based Agents. Understanding these foundational elements is key to building intelligent systems that can reason and act effectively in complex real-world scenarios.



# Discrete vs. Continuous Environments: The AI Playgrounds



## Discrete Environments

Imagine a game of chess. Each move is a distinct action, and the board state changes in clear, quantifiable steps. This is a discrete environment. In such settings, there are a finite number of states and actions. The agent's perception is clear, and the outcomes of actions are often predictable. Think of turn-based games, simple robotic tasks in a factory, or automated traffic light systems.



## Continuous Environments

Now, picture a self-driving car navigating a bustling city. The car's speed, steering angle, and the positions of other vehicles are constantly changing. These are continuous variables. In a continuous environment, there are infinitely many states and actions, making precise measurement and prediction much harder. Real-world scenarios like robotics, autonomous navigation, and natural language processing often fall into this category, requiring more sophisticated AI techniques.

Understanding the nature of the environment—whether it's discrete or continuous—is fundamental because it dictates the type of AI algorithms and representations an agent will need to employ. A discrete environment often allows for simpler, state-space search algorithms, while continuous environments typically demand more advanced techniques like reinforcement learning or neural networks that can handle continuous input and output.

# Characteristics of AI Environments: What Makes Them Unique?

## Fully Observable vs. Partially Observable

**Fully Observable:**Partially Observable:

## Deterministic vs. Stochastic

**Deterministic:**Stochastic:

## Episodic vs. Sequential

**Episodic:**Sequential:

## Static vs. Dynamic

**Static:**Dynamic:



These characteristics are critical for defining the complexity of the problem an AI agent faces. A fully observable, deterministic, episodic, and static environment is the simplest, while a partially observable, stochastic, sequential, and dynamic environment represents the pinnacle of AI challenges.

# Class Activity: Identifying Agents and Environments

Let's put our knowledge to the test! For each scenario below, identify the AI agent, its type, and describe the characteristics of its environment based on what we've learned.

1

## Scenario 1: Automated Vacuum Cleaner

- **Agent:** The vacuum cleaner itself.
- **Agent Type:** Simple Reflex Agent (responds to immediate perceptions like dirt) or Model-Based Reflex Agent (maintains an internal state of the room).
- **Environment Characteristics:**
  - Partially Observable:
  - Stochastic:
  - Sequential:
  - Dynamic:
  - Continuous:

2

## Scenario 2: Online Chess Game AI

- **Agent:** The chess AI program.
- **Agent Type:** Goal-Based Agent (aims to win the game) or Utility-Based Agent (maximizes its chances of winning).
- **Environment Characteristics:**
  - Fully Observable:
  - Deterministic:
  - Sequential:
  - Static:
  - Discrete:

3

## Scenario 3: Medical Diagnosis System

- **Agent:** The diagnostic AI software.
- **Agent Type:** Inference-Based Agent (uses rules and knowledge to deduce a diagnosis).
- **Environment Characteristics:**
  - Partially Observable:
  - Stochastic:
  - Episodic:
  - Static:
  - Discrete:

This exercise highlights how varied AI applications are and how tailoring the agent's design to its environment is crucial for success.

# Introduction to Knowledge-Based Agents (KBAs)

So far, we've discussed agents that react to direct perceptions or track environment states. Now, let's introduce a more sophisticated type: the **Knowledge-Based Agent (KBA)**. Unlike simpler agents, a KBA maintains an internal [knowledge base \(KB\)](#) – a collection of sentences representing facts about the world. It uses this knowledge to reason, make inferences, and decide on actions, rather than just reacting to sensory input. This allows KBAs to handle more complex tasks, solve problems, and even learn from experience by updating their knowledge base.

## Key Idea: "What it knows, not just what it sees"

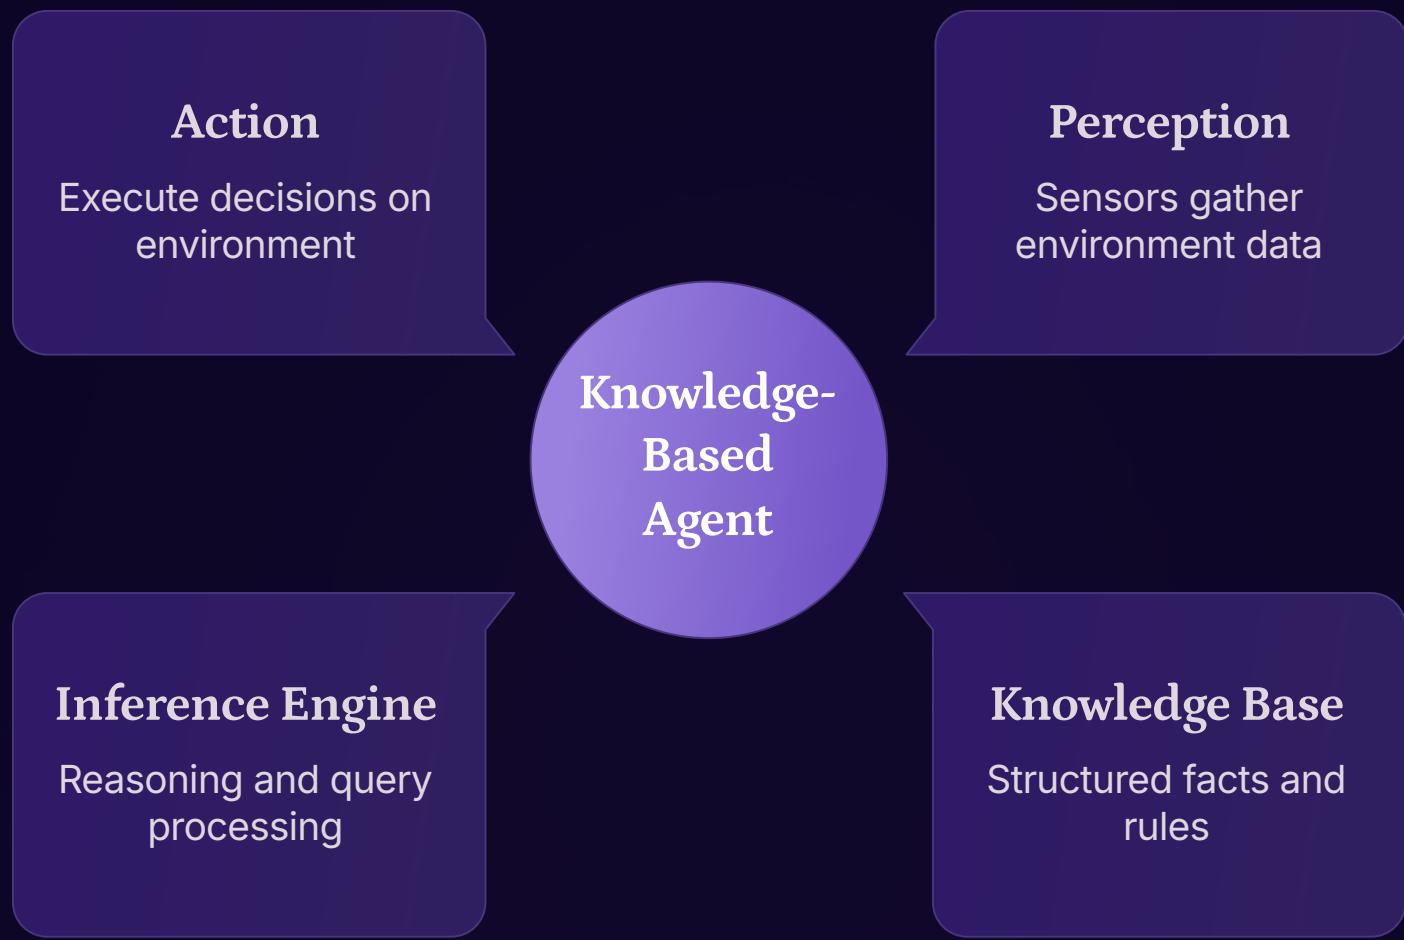
A KBA is essentially a reasoning system. It possesses a certain level of "understanding" about its domain. For instance, a simple vacuum cleaner might just react to seeing dirt. A KBA vacuum cleaner, however, might know the layout of the house, where dirt is most likely to accumulate, and even the optimal path to clean the entire area, not just the visible spots.

```
# Example: Simple Knowledge Base in Python (Conceptual)
knowledge_base = [
    "Mammals have fur.",
    "Dogs are mammals.",
    "Fido is a dog."
]

def ask_kba(query):
    if query == "Does Fido have fur?":
        # KBA would logically deduce this from its knowledge base
        return True
    return False

# Agent's action based on knowledge
if ask_kba("Does Fido have fur?"):
    print("Fido has fur!")
```

# Architecture of Knowledge-Based Agents



The architecture of a KBA is designed to facilitate reasoning and decision-making based on stored knowledge. It typically consists of two main components:

- **Knowledge Base (KB):**
- **Inference Engine:**

These two components work in tandem with the agent's perception and action mechanisms:

- **Perception:**
- **Tell:**
- **Ask:**
- **Action:**

This cyclical process allows the KBA to continuously perceive, reason, and act intelligently.

# Operations and Levels of KBA: From Sensing to Acting

## Operations of a KBA

A Knowledge-Based Agent performs a series of fundamental operations to achieve its goals:

1. **TELL:**
2. **ASK:**
3. **ACT:**

## Levels of KBA

To better understand and design KBAs, we can conceptualize them at different levels of abstraction:

01

### Knowledge Level

This is the most abstract level. It describes what the agent knows and what its goals are, without considering how that knowledge is represented or processed. At this level, we describe the agent's behavior purely in terms of its knowledge and beliefs about the world. For example, "The agent knows that opening the door will lead it to the next room."

02

### Logical Level

This level describes how knowledge is formally represented using a logical language (e.g., propositional logic, first-order logic) and how reasoning is performed using logical inference rules. Here, we specify the syntax and semantics of the knowledge base and the rules the inference engine uses. For example, representing "All humans are mortal" as  $\text{Human}(x) \Rightarrow \text{Mortal}(x)$ .

03

### Implementation Level

This is the concrete programming level. It details how the logical representations and inference mechanisms are actually coded in a programming language. This includes data structures for the knowledge base, algorithms for the inference engine, and the physical realization of the agent's perception and action systems. This level deals with the practicalities of making the KBA run on a computer.

# The Knowledge Level: What the Agent Knows

"At the knowledge level, an agent is described by what it knows, not by the symbols in which that knowledge is expressed." — Allen Newell

The Knowledge Level focuses on the agent's beliefs and objectives, treating the agent as if it were a black box that "thinks." We ask questions like: "What does the agent believe about the state of the world?" or "What goals is it trying to achieve?" This level is independent of the actual implementation details. It's about the **information content** of the agent's knowledge.

## Example: A Tourist Agent

At the Knowledge Level, we might say:

- The tourist agent **knows** that Paris is in France.
- It **knows** that the Eiffel Tower is a major landmark in Paris.
- It **desires** to visit major landmarks.
- It **believes** it has enough money for the trip.

We don't care *\*how\** it knows these facts (e.g., is it stored in a database, inferred from rules, or learned?). We only care that it possesses this knowledge and that this knowledge drives its rational actions (e.g., planning a trip to Paris to see the Eiffel Tower).

```
# Conceptual code: Knowledge Level perspective  
# Agent's knowledge might be represented internally,  
# but at this level, we only describe the knowledge itself.
```

```
# Agent knows:  
# location("Paris", "France")  
# landmark("Eiffel Tower", "Paris")  
# goal("visit", "landmark")  
# financial_status("sufficient_funds")
```

```
# Based on this, it acts rationally:  
# plan_trip("Paris")
```

# The Logical Level: How Knowledge is Represented and Used

The Logical Level takes the abstract knowledge from the Knowledge Level and gives it a concrete, formal structure. Here, we choose a specific representation language (e.g., propositional logic, first-order predicate logic, description logic) to encode the agent's knowledge. This language has a precise syntax (rules for forming valid sentences) and semantics (meaning of sentences).

Moreover, this level defines the inference rules that allow the agent to deduce new facts from existing ones. This is where the power of logical reasoning comes into play. For example, if the KB contains "All birds can fly" and "Tweety is a bird," the inference engine, using a rule like Modus Ponens, can deduce "Tweety can fly."

## Example: Formalizing the Tourist Agent's Knowledge

Using a simplified first-order logic, the tourist agent's knowledge might be represented as:

```
# Knowledge Base using a logical language (conceptual)

# Facts about location
Predicate: In(Location, Country)
Fact: In(Paris, France).

# Facts about landmarks
Predicate: IsLandmark(Landmark, Location)
Fact: IsLandmark(EiffelTower, Paris).

# Goal representation
Predicate: HasGoal(Agent, GoalType, Entity)
Fact: HasGoal(TouristAgent, Visit, landmark).

# Belief about financial status
Predicate: HasFunds(Agent, Status)
Fact: HasFunds(TouristAgent, Sufficient).

# Inference rule example:
# IF In(L, C) AND IsLandmark(LM, L) AND HasGoal(Agent, Visit, landmark)
# THEN AgentShouldVisit(LM, L)
```

At this level, we are concerned with the correctness and completeness of our logical representation and the soundness of our inference mechanisms. We want to ensure that our rules correctly derive true conclusions from true premises.

# The Implementation Level: Bringing AI to Life

Finally, the Implementation Level is where theory meets practice. This is the nuts and bolts of building the KBA. It involves choosing specific [data structures](#) to store the knowledge base (e.g., lists, databases, semantic networks, graph databases) and writing the [algorithms](#) for the inference engine (e.g., backward chaining, forward chaining, resolution). It's also where the sensory input is processed and translated into logical sentences (percepts), and the logical actions are translated into physical commands for actuators.

## Example: Python Implementation Snippets

For our tourist agent, the implementation might involve:

```
# Implementation Level: Python example

# Data structure for Knowledge Base
kb = {
    "facts": [
        {"predicate": "In", "args": ["Paris", "France"]},
        {"predicate": "IsLandmark", "args": ["EiffelTower", "Paris"]},
        {"predicate": "HasGoal", "args": ["TouristAgent", "Visit", "landmark"]},
        {"predicate": "HasFunds", "args": ["TouristAgent", "Sufficient"]}
    ],
    "rules": [
        # Example rule:
        # If agent has a goal to visit a landmark in a location,
        # and it has sufficient funds, it should plan a trip.
        {
            "antecedent": [
                {"predicate": "HasGoal", "args": ["Agent", "Visit", "landmark"]},
                {"predicate": "IsLandmark", "args": ["Landmark", "Location"]},
                {"predicate": "In", "args": ["Location", "Country"]},
                {"predicate": "HasFunds", "args": ["Agent", "Sufficient"]}
            ],
            "consequent": {"predicate": "Action", "args": ["Agent", "PlanTrip", "Location"]}
        }
    ]
}

# Simple Inference Engine (conceptual)
def infer_action(agent_kb):
    for rule in agent_kb["rules"]:
        # Check if all antecedents are met
        # (simplified: in a real system this would be a complex matching process)
        if all(fact_match(ant, agent_kb["facts"]) for ant in rule["antecedent"]):
            consequent = rule["consequent"]
            if consequent["predicate"] == "Action":
                return f"{consequent['args'][0]} should {consequent['args'][1]} to {consequent['args'][2]}"
    return "No action inferred"

# (Helper function for demonstration, would be more robust in real code)
def fact_match(query_fact, facts):
    for fact in facts:
        if fact["predicate"] == query_fact["predicate"] and \
           all(q_arg == f_arg or q_arg == "Agent" or q_arg == "Location" or q_arg == "Country" or q_arg == "Landmark" \
               for q_arg, f_arg in zip(query_fact["args"], fact["args"])):
            return True
    return False

# Example usage
print(infer_action(kb)) # Output: TouristAgent should PlanTrip to Paris (conceptual)
```

## Practice Questions for Students:

1. Describe the key difference between a discrete and a continuous environment, providing an example for each beyond those discussed.
2. Choose one characteristic of an environment (e.g., observability, determinism) and explain how it impacts the complexity of designing an AI agent.
3. What is the primary function of a Knowledge-Based Agent, and how does it differ from a simple reflex agent?
4. Imagine you are building an AI to play Tic-Tac-Toe. Describe how this agent would function at the Knowledge Level, Logical Level, and Implementation Level.
5. Consider a smart home assistant. Give an example of a "TELL" operation and an "ASK" operation it might perform with its knowledge base.